

SIMATS ENGINEERING (SAVEETHA SCHOOL OF ENGINEERING)

Department of Artificial Intelligence and Data Science

ASSESSMENT TOOL 3 – DESIGN TASK

Course Code:	DSA0306	Course Title:	Natural Language Processing
CO Assessed:	CO3 – Precision	Assessment:	Assessment Tool 3 – Design Task
Weightage:	40% of CIA		
CO5 Statement:	Apply N-gram models, smoothing techniques, part-of-speech tagging systems and to evaluate effective syntactic analysis. (BL-3)		
Student Name:	Y SRAVAN KUMAR	Reg. No.:	192472289
Section / Batch:	CSE AI 2024	Date:	18-08-2026

Code of Conduct

I Y SRAVAN KUMAR (192472289) (Name / Reg No) certify that this submission is my original work and that I have adhered to the guidelines specified for this assessment. I understand that any violation of academic integrity rules will result in disciplinary action.

Signature of the Student: _____ SRAVAN _____

Instructions

- Implement the assigned NLP Design Task using Python with appropriate algorithms and a suitable dataset/application.
- Execute the program and demonstrate the output using relevant sample inputs and test cases.
- Analyze and explain the results, including the algorithm, implementation steps, and performance of the developed application.
- Duration: 30 minutes.

Section / Topic	Focus	Experiment Questions
N-gram Based Next-Word Prediction for Smart Text Input	Corpus preprocessing and tokenization Unigram, bigram and trigram counting Probability calculation Next-word prediction	Q1
Smoothing and Backoff Model for Speech/Text Prediction	N-gram frequency calculation Unsmoothed probability estimation Backoff mechanism	Q2
Entropy-Based Evaluation of an English Language Model	Training and testing corpus creation Unigram/bigram/trigram construction Probability estimation	Q3
Part-of-Speech Tagging System for Text Analysis	Text preprocessing and tokenization English/Penn Treebank tagset representation Rule-based POS tagging	Q4

Task 1 – N-gram Based Text Prediction

Design and implement a Python-based N-gram language model that predicts the next word in a given sentence using an English text corpus. The system should preprocess the corpus, perform sentence segmentation and tokenization, and construct unigram, bigram, and trigram models by calculating word-frequency counts and maximum-likelihood probabilities.

For an incomplete sentence such as:

"Artificial intelligence is"

the system should identify and display the top-5 most probable next words.

The program should allow the user to select $N = 1, 2$, or 3 , display the corresponding N-gram frequency counts and probabilities, and demonstrate the behavior of the model when an unseen N-gram is encountered. Finally, test the model using multiple incomplete sentences, calculate the prediction accuracy, and discuss the limitations of using an unsmoothed N-gram model for real-world language prediction.

Task 2 – Backoff and Interpolation for Language Prediction

Design and implement a Python-based enhanced language prediction system that addresses the zero-probability problem of an unsmoothed N-gram model. The system should construct unigram, bigram, and trigram models from an English corpus and accept an incomplete sentence such as:

"Machine learning can"

When the required trigram is unavailable, the system should use a backoff strategy by first checking the trigram model, then the bigram model, and finally the unigram model. Extend the system using Deleted Interpolation to combine unigram, bigram, and trigram probabilities using predefined interpolation weights.

The system should generate the top-5 next-word predictions using:

- Unsmoothed N-gram model
- Backoff model
- Deleted Interpolation model

Compare the three approaches in terms of zero probabilities, prediction coverage, and prediction accuracy, and explain why backoff and interpolation are useful for sparse language data.

Task 3 – Entropy-Based Evaluation of Language Models

Design and implement a Python program for measuring the uncertainty of N-gram language models using entropy. Given separate training and testing corpora, construct unigram, bigram, and trigram language models and calculate the probabilities assigned to words in the test sentences. For each test sentence, calculate its entropy and classify the sentence as having relatively high or low predictive uncertainty.

For example, compare sentences such as:

"The cat sits on the mat."

and

"The quantum processor redesigned the..."

The system should evaluate how predictable the next word is based on the trained language model.

The program should also compare entropy values obtained using:

- Unigram model
- Bigram model
- Trigram model
- Smoothed model

Finally, interpret the results and explain how entropy can be used to measure prediction uncertainty and evaluate the reliability of a language model.

Task 4 – Comparative POS Tagging System

Design and implement a Python-based POS tagging system that assigns grammatical tags to words in English sentences.

The system should implement three approaches:

1. Rule-Based POS Tagging

Use a combination of lexical dictionaries, word patterns, and grammatical rules to assign POS tags.

2. Stochastic POS Tagging

Construct a statistical POS tagger using:

- Word/tag frequencies
- Tag transition probabilities
- Probabilities obtained from a training corpus

3. Transformation-Based Tagging

Initially assign tags using a simple tagging strategy and subsequently apply transformation rules to correct likely tagging errors.

For example:

"I book a ticket."

and

"I read a book."

The system should demonstrate how the same word can receive different POS tags depending on its context. Use an appropriate English corpus such as the Brown Corpus or another publicly available tagged corpus. The program should accept user-entered sentences, display the POS tags produced by all three approaches, and compare their results using suitable evaluation measures such as tagging accuracy. Finally, analyze the differences among the three approaches and explain the advantages and limitations of rule-based, stochastic, and transformation-based POS tagging.

Rubrics for Calculation

Criteria	Weightage (%)	Level 4 – Exemplary (4)	Level 3 – Proficient (3)	Level 2 – Developing (2)	Level 1 – Beginning (1)
Understanding of Concept	25%	Demonstrates complete understanding of design requirements and concepts	Good understanding with minor gaps	Basic understanding with some errors	Poor understanding or incorrect concepts
Design / Implementation	30%	Well-structured, efficient, and responsive design implementation	Correct implementation with minor issues	Basic implementation with inconsistencies	Incorrect or incomplete implementation
Use of Tools	20%	Effective and advanced use of tools/techniques	Appropriate use of tools	Limited or inconsistent use	Incorrect or minimal use
Analysis & Evaluation	15%	Insightful analysis with strong justification and optimization	Good analysis with justification	Basic analysis with limited depth	Poor or no analysis
Documentation / Clarity	10%	Clear, well-structured, and properly presented solution	Mostly clear with minor issues	Basic clarity, missing details	Poor or unclear documentation

Q. No. / Criteria Level	Understanding of Concepts	Experimental Design	Use of Tools	Analysis of Results	Documentation	Sub. Total	Total %
1	3	3	3	3	4	16	78
2	3	3	3	3	4	16	
3	3	3	3	3	4	16	
4	3	3	3	3	4	16	

Faculty In-charge	Course Coordinator	Program Director
Signature: _____ _DR.T.Kumaragurubaran_____	Signature: _____	Signature: _____
Date: _18-08-2026 _____	Date: _____	Date: _____

1.N-Gram Based Text Prediction Using a Language Model

CODE

```
nlp co3 at1 exp1.py - C:/Users/Sravan Kumar.Y/AppData/Local/Programs/Python/Python312/nlp co3 at1 exp1.py (3.12.3)
File Edit Format Run Options Window Help

import re
from collections import Counter, defaultdict

# =====
# 1. ENGLISH TEXT CORPUS
# =====

corpus = """
Artificial intelligence is changing the world.
Artificial intelligence is used in healthcare.
Artificial intelligence is transforming education.
Artificial intelligence is helping businesses.
Artificial intelligence is becoming important in modern technology.

Machine learning is a branch of artificial intelligence.
Machine learning is used to analyze data.
Machine learning is useful for prediction.
Machine learning can solve complex problems.

Natural language processing is a field of artificial intelligence.
Natural language processing is used to understand text.
Natural language processing helps computers understand language.

Deep learning is a part of machine learning.
Deep learning is useful for image recognition.
Deep learning is useful for speech recognition.

Data science is an important field of technology.
Data science is used for data analysis.
Data science helps organizations make decisions.

Computer science is the study of computation.
Computer science is important for modern technology.
Technology is changing the way people work.
Technology is improving our daily lives.
"""

# =====
# 2. SENTENCE SEGMENTATION
# =====

def sentence_segmentation(text):
    sentences = re.split(r'[\.\?!\,\;\:\-\&\+\=\>\<]+' text)
```

OUTPUT

```
IDLE Shell 3.12.3
File Edit Shell Debug Options Window Help
Python 3.12.3 (tags/v3.12.3:f6650f9, Apr 9 2024, 14:05:25) [MSC v.1938 64 bit (AMD64)] on win32
Type "help", "copyright", "credits" or "license()" for more information.
>>>
= RESTART: C:/Users/Sravan Kumar.Y/AppData/Local/Programs/Python/Python312/nlp co3 at1 expl.py
=====
PREPROCESSED CORPUS
=====
['<START>', 'artificial', 'intelligence', 'is', 'changing', 'the', 'world', '<END>']
['<START>', 'artificial', 'intelligence', 'is', 'used', 'in', 'healthcare', '<END>']
['<START>', 'artificial', 'intelligence', 'is', 'transforming', 'education', '<END>']
['<START>', 'artificial', 'intelligence', 'is', 'helping', 'businesses', '<END>']
['<START>', 'artificial', 'intelligence', 'is', 'becoming', 'important', 'in', 'modern', 'technology', '<END>']
['<START>', 'machine', 'learning', 'is', 'a', 'branch', 'of', 'artificial', 'intelligence', '<END>']
['<START>', 'machine', 'learning', 'is', 'used', 'to', 'analyze', 'data', '<END>']
['<START>', 'machine', 'learning', 'is', 'useful', 'for', 'prediction', '<END>']
['<START>', 'machine', 'learning', 'can', 'solve', 'complex', 'problems', '<END>']
['<START>', 'natural', 'language', 'processing', 'is', 'a', 'field', 'of', 'artificial', 'intelligence', '<END>']
['<START>', 'natural', 'language', 'processing', 'is', 'used', 'to', 'understand', 'text', '<END>']
['<START>', 'natural', 'language', 'processing', 'helps', 'computers', 'understand', 'language', '<END>']
['<START>', 'deep', 'learning', 'is', 'a', 'part', 'of', 'machine', 'learning', '<END>']
['<START>', 'deep', 'learning', 'is', 'useful', 'for', 'image', 'recognition', '<END>']
['<START>', 'deep', 'learning', 'is', 'useful', 'for', 'speech', 'recognition', '<END>']
['<START>', 'data', 'science', 'is', 'an', 'important', 'field', 'of', 'technology', '<END>']
['<START>', 'data', 'science', 'is', 'used', 'for', 'data', 'analysis', '<END>']
['<START>', 'data', 'science', 'helps', 'organizations', 'make', 'decisions', '<END>']
['<START>', 'computer', 'science', 'is', 'the', 'study', 'of', 'computation', '<END>']
['<START>', 'computer', 'science', 'is', 'important', 'for', 'modern', 'technology', '<END>']
['<START>', 'technology', 'is', 'changing', 'the', 'way', 'people', 'work', '<END>']
['<START>', 'technology', 'is', 'improving', 'our', 'daily', 'lives', '<END>']

=====
N-GRAM LANGUAGE MODEL
=====

Select N:
1 - Unigram
2 - Bigram
3 - Trigram

Enter N (1, 2, or 3): 2

=====
2-GRAM FREQUENCY COUNTS AND PROBABILITIES
=====
('artificial', 'intelligence') Count = 7 Probability = 1.0000
```

CODE

```
import re
from collections import Counter, defaultdict
```

```
# =====
# 1. ENGLISH TEXT CORPUS
# =====
```

```
corpus = """
Artificial intelligence is changing the world.
Artificial intelligence is used in healthcare.
Artificial intelligence is transforming education.
Artificial intelligence is helping businesses.
Artificial intelligence is becoming important in modern technology.
```

```
Machine learning is a branch of artificial intelligence.
Machine learning is used to analyze data.
Machine learning is useful for prediction.
Machine learning can solve complex problems.
```

```
Natural language processing is a field of artificial intelligence.
Natural language processing is used to understand text.
Natural language processing helps computers understand language.
```

Deep learning is a part of machine learning.
Deep learning is useful for image recognition.
Deep learning is useful for speech recognition.

Data science is an important field of technology.
Data science is used for data analysis.
Data science helps organizations make decisions.

Computer science is the study of computation.
Computer science is important for modern technology.
Technology is changing the way people work.
Technology is improving our daily lives.

"""

```
# =====  
# 2. SENTENCE SEGMENTATION  
# =====
```

```
def sentence_segmentation(text):  
    sentences = re.split(r'[.!?]+' , text)  
  
    sentences = [  
        sentence.strip()  
        for sentence in sentences  
        if sentence.strip()  
    ]  
  
    return sentences
```

```
# =====  
# 3. TOKENIZATION  
# =====
```

```
def tokenize(sentence):  
    return re.findall(r'\b[a-zA-Z]+\b' , sentence.lower())
```

```
# =====  
# 4. PREPROCESS THE CORPUS  
# =====
```

```
sentences = sentence_segmentation(corpus)
```

```
tokenized_sentences = []
```

```
for sentence in sentences:  
    tokens = tokenize(sentence)
```

```
    # Add sentence boundary tokens  
    tokens = ["<START>"] + tokens + ["<END>"]
```

```

tokenized_sentences.append(tokens)

print("=" * 60)
print("PREPROCESSED CORPUS")
print("=" * 60)

for sentence in tokenized_sentences:
    print(sentence)

# =====
# 5. BUILD UNIGRAM, BIGRAM AND TRIGRAM COUNTS
# =====

unigram_counts = Counter()
bigram_counts = Counter()
trigram_counts = Counter()

for tokens in tokenized_sentences:

    # Unigram
    for word in tokens:
        unigram_counts[word] += 1

    # Bigram
    for i in range(len(tokens) - 1):
        bigram = (tokens[i], tokens[i + 1])
        bigram_counts[bigram] += 1

    # Trigram
    for i in range(len(tokens) - 2):
        trigram = (tokens[i], tokens[i + 1], tokens[i + 2])
        trigram_counts[trigram] += 1

# =====
# 6. MAXIMUM LIKELIHOOD PROBABILITIES
# =====

# Unigram probability:
#  $P(\text{word}) = \text{Count}(\text{word}) / \text{Total words}$ 

total_words = sum(unigram_counts.values())

unigram_probabilities = {}

for word, count in unigram_counts.items():
    unigram_probabilities[word] = count / total_words

# Bigram probability:

```

```

#  $P(\text{word2} \mid \text{word1}) =$ 
#  $\text{Count}(\text{word1}, \text{word2}) / \text{Count}(\text{word1})$ 

bigram_probabilities = {}

for (word1, word2), count in bigram_counts.items():
    bigram_probabilities[(word1, word2)] = (
        count / unigram_counts[word1]
    )

# Trigram probability:
#  $P(\text{word3} \mid \text{word1}, \text{word2}) =$ 
#  $\text{Count}(\text{word1}, \text{word2}, \text{word3}) / \text{Count}(\text{word1}, \text{word2})$ 

trigram_probabilities = {}

for (word1, word2, word3), count in trigram_counts.items():
    trigram_probabilities[(word1, word2, word3)] = (
        count / bigram_counts[(word1, word2)]
    )

# =====
# 7. DISPLAY N-GRAM COUNTS AND PROBABILITIES
# =====

def display_ngrams(n):

    print("\n" + "=" * 60)
    print(f"{n}-GRAM FREQUENCY COUNTS AND PROBABILITIES")
    print("=" * 60)

    if n == 1:

        for word, count in unigram_counts.most_common():
            probability = unigram_probabilities[word]

            print(
                f"{word:15s} Count = {count:3d} "
                f"Probability = {probability:.4f} "
            )

    elif n == 2:

        for bigram, count in bigram_counts.most_common():
            probability = bigram_probabilities[bigram]

            print(
                f"{bigram} Count = {count:3d} "
                f"Probability = {probability:.4f} "
            )

```

```
elif n == 3:
```

```
    for trigram, count in trigram_counts.most_common():
        probability = trigram_probabilities[trigram]
```

```
    print(
        f'{trigram} Count = {count:3d} '
        f'Probability = {probability:.4f}'
    )
```

```
# =====
# 8. PREDICT NEXT WORD
# =====
```

```
def predict_next_word(sentence, n, top_k=5):
```

```
    tokens = tokenize(sentence)
```

```
    if len(tokens) == 0:
        return []
```

```
    predictions = []
```

```
# -----
# UNIGRAM
# -----
```

```
if n == 1:
```

```
    for word, probability in unigram_probabilities.items():
```

```
        if word not in ["<START>", "<END>"]:
            predictions.append((word, probability))
```

```
# -----
# BIGRAM
# -----
```

```
elif n == 2:
```

```
    last_word = tokens[-1]
```

```
    for (word1, word2), probability in bigram_probabilities.items():
```

```
        if word1 == last_word and word2 not in ["<START>", "<END>"]:
            predictions.append((word2, probability))
```

```
# -----
# TRIGRAM
# -----
```

```
elif n == 3:
```

```

    if len(tokens) < 2:
        return []

    word1 = tokens[-2]
    word2 = tokens[-1]

    for (w1, w2, w3), probability in trigram_probabilities.items():

        if w1 == word1 and w2 == word2:
            if w3 not in ["<START>", "<END>"]:
                predictions.append((w3, probability))

    predictions.sort(key=lambda x: x[1], reverse=True)

    return predictions[:top_k]

# =====
# 9. DISPLAY TOP-5 PREDICTIONS
# =====

def display_prediction(sentence, n):

    print("\n" + "=" * 60)
    print("NEXT WORD PREDICTION")
    print("=" * 60)

    print("Sentence:", sentence)
    print("N =", n)

    predictions = predict_next_word(sentence, n)

    if not predictions:

        print("\nNo prediction available.")
        print("This is an unseen N-gram.")

        return

    print("\nTop-5 predicted next words:")

    for rank, (word, probability) in enumerate(predictions, 1):

        print(
            f"{rank}. {word:<15} "
            f"Probability = {probability:.4f}"
        )

# =====
# 10. UNSEEN N-GRAM DEMONSTRATION
# =====

```

```

def check_unseen_ngram(sentence, n):

    tokens = tokenize(sentence)

    print("\n" + "=" * 60)
    print("UNSEEN N-GRAM CHECK")
    print("=" * 60)

    if n == 2:

        if len(tokens) < 2:
            return

        bigram = (tokens[-2], tokens[-1])

        if bigram not in bigram_counts:

            print("Unseen Bigram:", bigram)
            print("Frequency = 0")
            print("Probability = 0")

        else:

            print("Bigram exists:", bigram)
            print(
                "Probability =",
                bigram_probabilities[bigram]
            )

    elif n == 3:

        if len(tokens) < 3:
            return

        trigram = (
            tokens[-3],
            tokens[-2],
            tokens[-1]
        )

        if trigram not in trigram_counts:

            print("Unseen Trigram:", trigram)
            print("Frequency = 0")
            print("Probability = 0")

        else:

            print("Trigram exists:", trigram)
            print(
                "Probability =",
                trigram_probabilities[trigram]
            )

```

)

```
# =====  
# 11. ACCURACY TEST  
# =====
```

```
test_cases = [  
    ("Artificial intelligence is", "changing"),  
    ("Machine learning is", "a"),  
    ("Natural language processing is", "a"),  
    ("Deep learning is", "a"),  
    ("Data science is", "an"),  
    ("Computer science is", "the"),  
]
```

```
def calculate_accuracy(n):  
    correct = 0  
    total = len(test_cases)  
  
    print("\n" + "=" * 60)  
    print(f"ACCURACY TEST - N = {n}")  
    print("=" * 60)  
  
    for sentence, actual_word in test_cases:  
        predictions = predict_next_word(  
            sentence,  
            n,  
            top_k=5  
        )  
  
        predicted_words = [  
            word for word, probability in predictions  
        ]  
  
        if actual_word in predicted_words:  
            correct += 1  
            result = "CORRECT"  
        else:  
            result = "WRONG"
```

```

print(
    f"\nSentence: {sentence}"
)

print(
    f"Actual next word: {actual_word}"
)

print(
    f"Predictions: {predicted_words}"
)

print(
    f"Result: {result}"
)

```

```

accuracy = (correct / total) * 100

```

```

print("\n" + "-" * 60)
print(
    f"Accuracy = {correct}/{total} "
    f"= {accuracy:.2f}%"
)
print("-" * 60)

```

```

# =====
# 12. MAIN PROGRAM
# =====

```

```

print("\n" + "=" * 60)
print("N-GRAM LANGUAGE MODEL")
print("=" * 60)

```

```

print("\nSelect N:")
print("1 - Unigram")
print("2 - Bigram")
print("3 - Trigram")

```

```

n = int(input("\nEnter N (1, 2, or 3): "))

```

```

if n not in [1, 2, 3]:

```

```

    print("Invalid choice.")

```

```

else:

```

```

    # Display frequency counts and probabilities
    display_ngrams(n)

```

```

    # Input sentence
    sentence = input(

```

```
"\nEnter an incomplete sentence: "
)
```

```
# Prediction
display_prediction(sentence, n)
```

```
# Unseen N-gram check
if n in [2, 3]:
    check_unseen_ngram(sentence, n)
```

```
# Accuracy
calculate_accuracy(n)
```

```
# =====
# 13. LIMITATIONS
# =====
```

```
print("\n" + "=" * 60)
print("LIMITATIONS OF UNSMOOTHED N-GRAM MODEL")
print("=" * 60)
```

```
print("""
1. Unseen N-grams get probability 0.

2. The model cannot predict word combinations that
   were not present in the training corpus.

3. It requires a large corpus for better predictions.

4. Unigram ignores word context.

5. Bigram and trigram models consider only limited context.

6. The model cannot understand the actual meaning of
   words or sentences.

7. The model may perform poorly on new or rare sentences.
""")
```